

Firts-year College GPA Prediction Using Gaussian Linear Model

Amanda Aarnio, Anni Niskanen, Antti Huttunen

2022-12-02

Contents

1	Introduction	2
2	Data and Problem	2
3	Models	2
4	Analysis and Results	8
5	Discussion about problems and improvements	20
6	Conclusion	20
7	Self-reflection	20

1 Introduction

Can the grade point averages (GPAs) of first-year college students be predicted based on high school GPAs, SAT scores and other factors? This report presents a solution to this problem by applying two statistical models to predict the first-year college GPA. These models are fitted in Stan and compared according to Bayesian data analysis concepts.

The report begins with a description of the data and the problem. After this, the models are introduced and their performances are analysed. Model improvements are then discussed and conclusions are made. Finally, self-reflection of the project is presented.

2 Data and Problem

The data contains information from a sample of 219 first year students at a Midwestern college in 1996. The data has 10 variables:

GPA	First-year college GPA on a 0.0 to 4.0 scale
HSGPA	High school GPA on a 0.0 to 4.0 scale
SATV	Verbal/critical reading SAT score
SATM	Math SAT score
Male	1= male, 0= female
HU	Number of credit hours earned in humanities courses in high school
SS	Number of credit hours earned in social science courses in high school
FirstGen	1= student is the first in her or his family to attend college, 0=otherwise
White	1= white students, 0= others
CollegeBound	1=attended a high school where >=50% students intended to go on to college, 0=otherwise

The description of the data can be obtained from: <https://vincentarelbundock.github.io/Rdatasets/doc/Stat2Data/FirstYearGPA.html>. The data can be accessed in R from the Stat2Data package as follows:

```
library(Stat2Data)
data("FirstYearGPA")
```

This report attempts to solve the problem of predicting the GPA of first-year college students using this data. Five variables were chosen to predict the first-year college GPA (GPA): HSGPA, SATV, SATM, HU and SS. All of the chosen variables have numerical values. In addition, the hierarchical model uses categorical variables Male and White to group the data into four different groups and predicts the GPA with group-level predictors. The two categorical variables were selected due to their most similar sample sizes. The barplot in the Figure 1 represents these four groups and their sizes.

```
# Importing packages
library(cmdstanr)
library(bayesplot)
library(ggplot2)
library(gridExtra)
library(loo)
```

3 Models

In this project, two different models are utilised: a pooled and a hierarchical model. Both are covered in following sections. The mathematical notations, Stan implementations, and Stan model runs are included in the sections. Both models are linear Gaussian models whose expected values μ_i are constructed as linear




Figure 1: The division of the data in four different groups. The amounts of the data points in the groups can be seen in bars.

combinations of the chosen variables HSGPA, SATV, SATM, HU, and SS. The coefficients of these linear combinations are parameters α , β_1 , β_2 , β_3 , β_4 , and β_5 .

Weakly informative normal priors were used in both models. $N(0,10)$ was the prior distribution for σ in both models. In the pooled model, $N(0,100)$ was the prior for both α and all β_k . The hierarchical model had similar priors: $N(0,100)$ for all means μ_α and μ_{β_k} , and $N(0,10)$ for all standard deviations σ_α and σ_{β_k} . These priors were considered to be reasonable for multiple reasons; Firstly, they all center on 0, which was thought wisest as there is no information whether the intercept α or the slopes β_k should be positive or negative. Secondly, the standard deviations, 100 for the means and 10 for the standard deviations, were considered large enough to produce wide enough (but not too wide) prior distributions.

As it can be seen from the code where the Stan models are run, default values were used for running the MCMC chains. 4 chains with 2000 iterations were run, and the first 1000 iterations of each chain were considered warm-up.

3.1 Pooled model

In the pooled model, the expected values μ_i are constructed using the common parameters α and β_k , which have the above-mentioned weakly informative priors. The GPAs have a common standard deviation σ .

Mathematical notation

$$\begin{aligned}
 GPA_i &\sim N(\mu_i, \sigma) \\
 \mu_i &= \alpha + \beta_1 \cdot HSGPA_i + \beta_2 \cdot SATV_i + \beta_3 \cdot SATM_i + \beta_4 \cdot HU_i + \beta_5 \cdot SS_i \\
 \sigma &\sim N(0, 10) \\
 \alpha &\sim N(0, 100) \\
 \beta_k &\sim N(0, 100)
 \end{aligned}$$

Stan code

```
data {
  int<lower=0> N;
```

```

matrix[N,5] x;
vector[N] y;
real musigma;
real sigmasigma;
}

parameters {
  real alpha;
  vector[5] betas;
  real<lower=0> sigma;
}

transformed parameters {
  vector[N] mu;
  mu = alpha + betas[1]*x[,1] + betas[2]*x[,2] + betas[3]*x[,3] +
        betas[4]*x[,4] + betas[5]*x[,5];
}

model {
  // priors
  alpha ~ normal(0, musigma);
  betas ~ normal(0, musigma);
  sigma ~ normal(0, sigmasigma);

  // likelihood
  y ~ normal(mu, sigma);
}

generated quantities {
  vector[N] ypred;
  vector[N] log_lik;

  // Generate predictive distributions for GPA
  for (i in 1:N)
    ypred[i] = normal_rng(mu[i], sigma);

  // log likelihoods
  for (i in 1:N)
    log_lik[i] = normal_lpdf(y[i] | mu[i], sigma);
}

```

Running the model

```

data_pooled <- list(N = nrow(FirstYearGPA),
  x = subset(FirstYearGPA, select = c('HSGPA', 'SATV', 'SATM',
                                     'HU', 'SS')),
  y = FirstYearGPA$GPA,
  musigma = 100,
  sigmasigma = 10)

```

```

mod_pooled <- cmdstan_model("pooled.stan")
fit_pooled <- mod_pooled$sample(data_pooled, refresh = 0, seed = 03091900)

```

```

## Running MCMC with 4 sequential chains...
##

```

```

## Chain 1 finished in 16.5 seconds.
## Chain 2 finished in 15.8 seconds.
## Chain 3 finished in 18.0 seconds.
## Chain 4 finished in 16.0 seconds.
##
## All 4 chains finished successfully.
## Mean chain execution time: 16.6 seconds.
## Total execution time: 67.0 seconds.

```

3.2 Hierarchical model

In the hierarchical model, students have been divided in four different groups: 1. white males, 2. non-white males, 3. white females, and 4. non-white females. The groups have their own parameters α_j and β_{kj} , where j denotes the group a student belongs to. The parameters α_j and β_{kj} have common hyperparameters, implying the parameters can be different across groups, but they have something common as well.

Mathematical notation

$$\begin{aligned}
GPA_{ij} &\sim N(\mu_{ij}, \sigma) \\
\mu_{ij} &= \alpha_j + \beta_{1j} \cdot HSGPA_i + \beta_{2j} \cdot SATV_i + \beta_{3j} \cdot SATM_i + \beta_{4j} \cdot HU_i + \beta_{5j} \cdot SS_i \\
\sigma &\sim N(0, 10) \\
\alpha_j &\sim N(\mu_\alpha, \sigma_\alpha) \\
\beta_{1j} &\sim N(\mu_{\beta_1}, \sigma_{\beta_1}) \\
\beta_{2j} &\sim N(\mu_{\beta_2}, \sigma_{\beta_2}) \\
\beta_{3j} &\sim N(\mu_{\beta_3}, \sigma_{\beta_3}) \\
\beta_{4j} &\sim N(\mu_{\beta_4}, \sigma_{\beta_4}) \\
\beta_{5j} &\sim N(\mu_{\beta_5}, \sigma_{\beta_5}) \\
\mu_\alpha &\sim N(0, 100) \\
\sigma_\alpha &\sim N(0, 10) \\
\mu_{\beta_k} &\sim N(0, 100) \\
\sigma_{\beta_k} &\sim N(0, 10)
\end{aligned}$$

Stan code

```

data {
  int<lower=0> N1;
  int<lower=0> N2;
  int<lower=0> N3;
  int<lower=0> N4;
  matrix[N1,5] x1;
  matrix[N2,5] x2;
  matrix[N3,5] x3;
  matrix[N4,5] x4;
  vector[N1] y1;
  vector[N2] y2;
  vector[N3] y3;
  vector[N4] y4;
  real musigma;
  real sigmasigma;
}

parameters {

```

```

// parameters
real alpha1;
real alpha2;
real alpha3;
real alpha4;
vector[5] betas1;
vector[5] betas2;
vector[5] betas3;
vector[5] betas4;
real<lower=0> sigma;

// hyperparameters
real pmualpha;
real<lower=0> psalpha;
vector[5] pmubetas;
vector<lower=0>[5] psbetas;
}

transformed parameters {
  vector[N1] mu1 = alpha1 + betas1[1]*x1[,1] + betas1[2]*x1[,2]
    + betas1[3]*x1[,3] + betas1[4]*x1[,4] + betas1[5]*x1[,5];
  vector[N2] mu2 = alpha2 + betas2[1]*x2[,1] + betas2[2]*x2[,2]
    + betas2[3]*x2[,3] + betas2[4]*x2[,4] + betas2[5]*x2[,5];
  vector[N3] mu3 = alpha3 + betas3[1]*x3[,1] + betas3[2]*x3[,2]
    + betas3[3]*x3[,3] + betas3[4]*x3[,4] + betas3[5]*x3[,5];
  vector[N4] mu4 = alpha4 + betas4[1]*x4[,1] + betas4[2]*x4[,2]
    + betas4[3]*x4[,3] + betas4[4]*x4[,4] + betas4[5]*x4[,5];
}

model {
  // hyperpriors
  pmualpha ~ normal(0, musigma);
  psalpha ~ normal(0, sigmasigma);
  for (i in 1:5){
    pmubetas[i] ~ normal(0, musigma);
    psbetas[i] ~ normal(0, sigmasigma);
  }

  // priors
  alpha1 ~ normal(pmualpha, psalpha);
  alpha2 ~ normal(pmualpha, psalpha);
  alpha3 ~ normal(pmualpha, psalpha);
  alpha4 ~ normal(pmualpha, psalpha);
  betas1 ~ normal(pmubetas, psbetas);
  betas2 ~ normal(pmubetas, psbetas);
  betas3 ~ normal(pmubetas, psbetas);
  betas4 ~ normal(pmubetas, psbetas);
  sigma ~ normal(0, sigmasigma);

  // likelihoods
  y1 ~ normal(mu1, sigma);
  y2 ~ normal(mu2, sigma);
  y3 ~ normal(mu3, sigma);

```

```

y4 ~ normal(mu4, sigma);
}

generated quantities{
  vector[N1] ypred1;
  vector[N2] ypred2;
  vector[N3] ypred3;
  vector[N4] ypred4;
  vector[N1+N2+N3+N4] log_lik;

  // Generate predictive distributions for GPA
  for (i in 1:N1)
    ypred1[i] = normal_rng(mu1[i], sigma);
  for (i in 1:N2)
    ypred2[i] = normal_rng(mu2[i], sigma);
  for (i in 1:N3)
    ypred3[i] = normal_rng(mu3[i], sigma);
  for (i in 1:N4)
    ypred4[i] = normal_rng(mu4[i], sigma);

  // log likelihoods
  for (i in 1:N1)
    log_lik[i] = normal_lpdf(y1[i] | mu1[i], sigma);
  for (i in 1:N2)
    log_lik[N1+i] = normal_lpdf(y2[i] | mu2[i], sigma);
  for (i in 1:N3)
    log_lik[N1+N2+i] = normal_lpdf(y3[i] | mu3[i], sigma);
  for (i in 1:N4)
    log_lik[N1+N2+N3+i] = normal_lpdf(y4[i] | mu4[i], sigma);
}

```

Running the model

```

male_white <- FirstYearGPA[FirstYearGPA$Male==1 & FirstYearGPA$White==1,]
male_non_white <- FirstYearGPA[FirstYearGPA$Male==1 & FirstYearGPA$White==0,]
female_white <- FirstYearGPA[FirstYearGPA$Male==0 & FirstYearGPA$White==1,]
female_non_white <- FirstYearGPA[FirstYearGPA$Male==0 & FirstYearGPA$White==0,]

data_hierarchical <- list(N1 = nrow(male_white),
                          N2 = nrow(male_non_white),
                          N3 = nrow(female_white),
                          N4 = nrow(female_non_white),
                          x1 = subset(male_white, select = c('HSGPA', 'SATV',
                                                            'SATM', 'HU',
                                                            'SS')),
                          x2 = subset(male_non_white, select = c('HSGPA', 'SATV',
                                                                'SATM', 'HU',
                                                                'SS')),
                          x3 = subset(female_white, select = c('HSGPA', 'SATV',
                                                                'SATM', 'HU',
                                                                'SS')),
                          x4 = subset(female_non_white, select = c('HSGPA', 'SATV',
                                                                    'SATM', 'HU',
                                                                    'SS'))),

```

```

      y1 = male_white$GPA,
      y2 = male_non_white$GPA,
      y3 = female_white$GPA,
      y4 = female_non_white$GPA,
      musigma = 100,
      sigmasigma = 10)

mod_hierarchical <- cmdstan_model("hierarchical.stan")
fit_hierarchical <- mod_hierarchical$sample(data_hierarchical, refresh = 0,
                                             seed = 03091900)

## Running MCMC with 4 sequential chains...
##
## Chain 1 finished in 31.7 seconds.
## Chain 2 finished in 29.7 seconds.
## Chain 3 finished in 34.1 seconds.
## Chain 4 finished in 35.7 seconds.
##
## All 4 chains finished successfully.
## Mean chain execution time: 32.8 seconds.
## Total execution time: 131.6 seconds.

## Warning: 218 of 4000 (5.0%) transitions ended with a divergence.
## See https://mc-stan.org/misc/warnings for details.

## Warning: 4 of 4000 (0.0%) transitions hit the maximum treedepth limit of 10.
## See https://mc-stan.org/misc/warnings for details.

```

4 Analysis and Results

This section contains the analysis of both models. The analysis includes convergence diagnostics, posterior predictive checks, and prior sensitivity analysis. This section also contains discussion about posterior distributions of the parameters and model comparison.

4.1 Converge diagnostics

4.1.1 Pooled model

```

summary_pooled <- fit_pooled$summary()
sum(summary_pooled$rhat > 1.01)
sum(summary_pooled$rhat < 0.99)

## [1] 0
## [1] 0

mean(summary_pooled$ess_tail)
sum(summary_pooled$ess_tail<400)
mean(summary_pooled$ess_bulk)
sum(summary_pooled$ess_bulk<400)

## [1] 3171.745
## [1] 0
## [1] 3207.508
## [1] 0

```


None of the \hat{R} -values for the pooled model exceed 1.01, which indicates that there are no problems with the convergence of the simulated chains. The effective sample size (ESS) is over 400 (recommended threshold with four parallel chains, 100 per chain) for all of the `ess_bulk` and `ess_tail` values, which indicates that the \hat{R} -estimate can be trusted to make decisions about convergence and the quality of the chains. Finally, the model does not observe divergence.

4.1.2 Hierarchical model

```
summary_hierarchical <- fit_hierarchical$summary()
sum(summary_hierarchical$rhat > 1.01)
sum(summary_hierarchical$rhat < 0.99)

## [1] 0
## [1] 0

mean(summary_hierarchical$ess_tail)
sum(summary_hierarchical$ess_tail<400)
mean(summary_hierarchical$ess_bulk)
sum(summary_hierarchical$ess_bulk<400)

## [1] 2984.828
## [1] 2
## [1] 2725.984
## [1] 4
```

Similarly to the pooled model, in the hierarchical model, no \hat{R} -values exceed 1.01, indicating that the chains converge. A few ESS values are below 400 and the model gives a divergence of approximately 5%. These facts indicate that there might be some problems with the convergence of the model. However, the obtained values do not differ significantly from the desired values, so the probability of having problems with convergence is small.

4.2 Posterior predictive checks

48 histograms of replications of the original data, i.e. the posterior predictive distributions, are shown for all models. In addition, the empirical probability density (epdf) and empirical cumulative distribution (ecdf) functions of the original data and 100 replications are overlapped and compared in order to investigate how well the posterior corresponds to the original data. The closer the replications are to the original data, the better predictions the fitted model can make.

4.2.1 Pooled model

As can be seen in Figure 2, the replicated histograms have a similar distribution and width compared to the original data. Moreover, the replicated epdf and ecdf functions follow the shape of the epdf and ecdf of the original data, and variation around the original functions is quite low. Therefore the model posterior corresponds to the original data well and the model is good in posterior prediction.

```
y <- FirstYearGPA$GPA
ypred <- fit_pooled$draws("ypred", format = "matrix")

grid.arrange( ppc_hist(y, ypred[1:48,]),
              ppc_dens_overlay(y, ypred[1:100,]),
              ppc_ecdf_overlay(y, ypred[1:100,]),
              layout_matrix = matrix(c(1,2,1,3), nrow = 2), nrow=2)
```

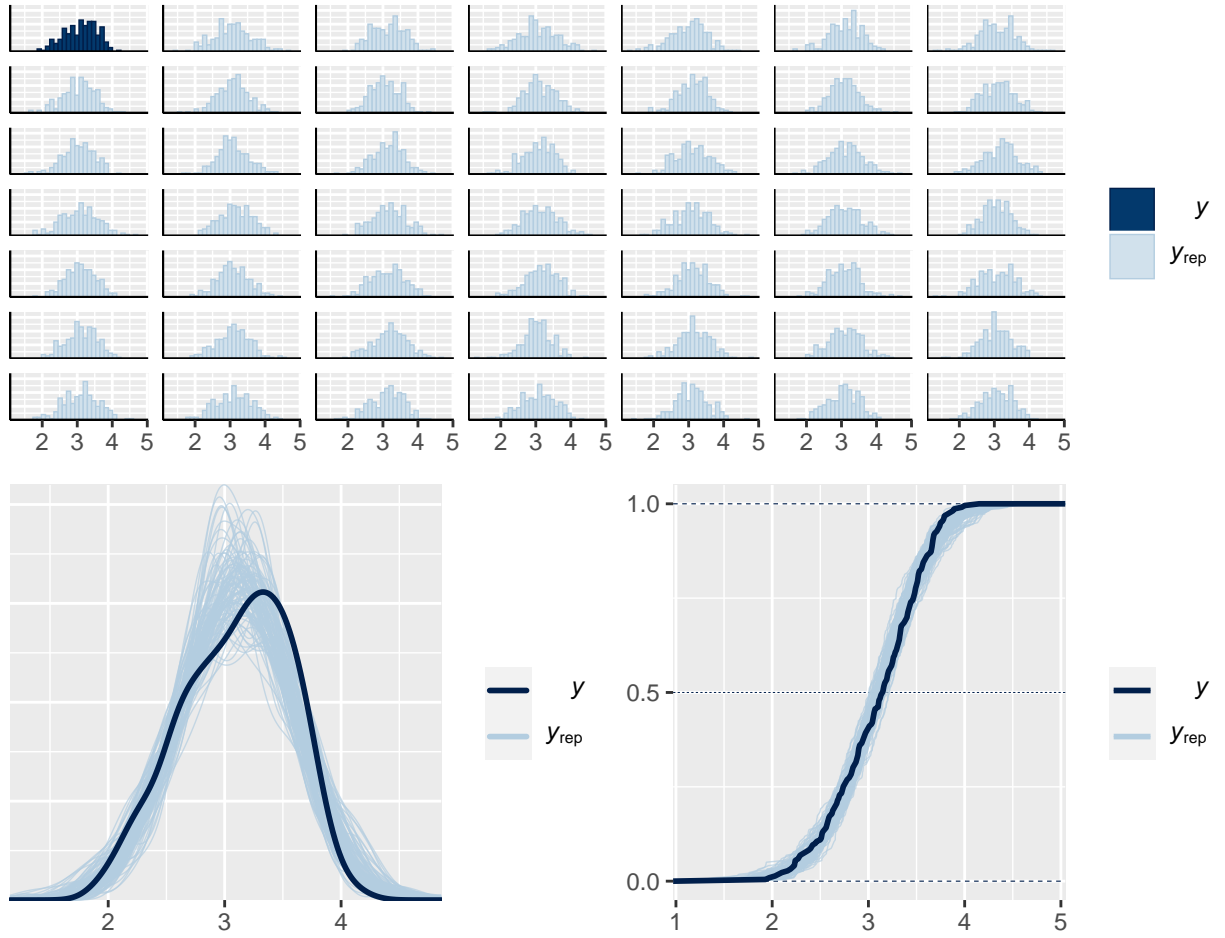


Figure 2: On the top of the figure 48 replications of the pooled model are shown with histograms. The original data is shown on the top left corner. Below the histograms, 100 epdf and ecdf distributions of the replications are overlapped with the epdf and ecdf of original data.

4.2.2 Hierarchical model

As can be seen when comparing Figures 3, 4, 5, and 6, the replications follow the original data better for the group of the white students than the non-white students. The biggest reason for this is the smaller amount of data points in non-white groups. It can be seen from the figures that replications of non-white groups have more width variation and more outliers than in the original data. In addition, ecdfs have more variation in the non-white groups than in the white groups.

Overall, the replications of the white groups are quite good, although not as good as in the pooled model, and the replications in non-white groups are not as good since they don't follow the original data as well.

```
y1 <- data_hierarchical$y1
ypred1 <- fit_hierarchical$draws("ypred1", format = "matrix")

grid.arrange(ppc_hist(y1, ypred1[1:48,]),
ppc_dens_overlay(y1, ypred1[1:100,]),
ppc_ecdf_overlay(y1, ypred1[1:100,]),
              layout_matrix = matrix(c(1,2,1,3), nrow = 2), nrow=2)
```

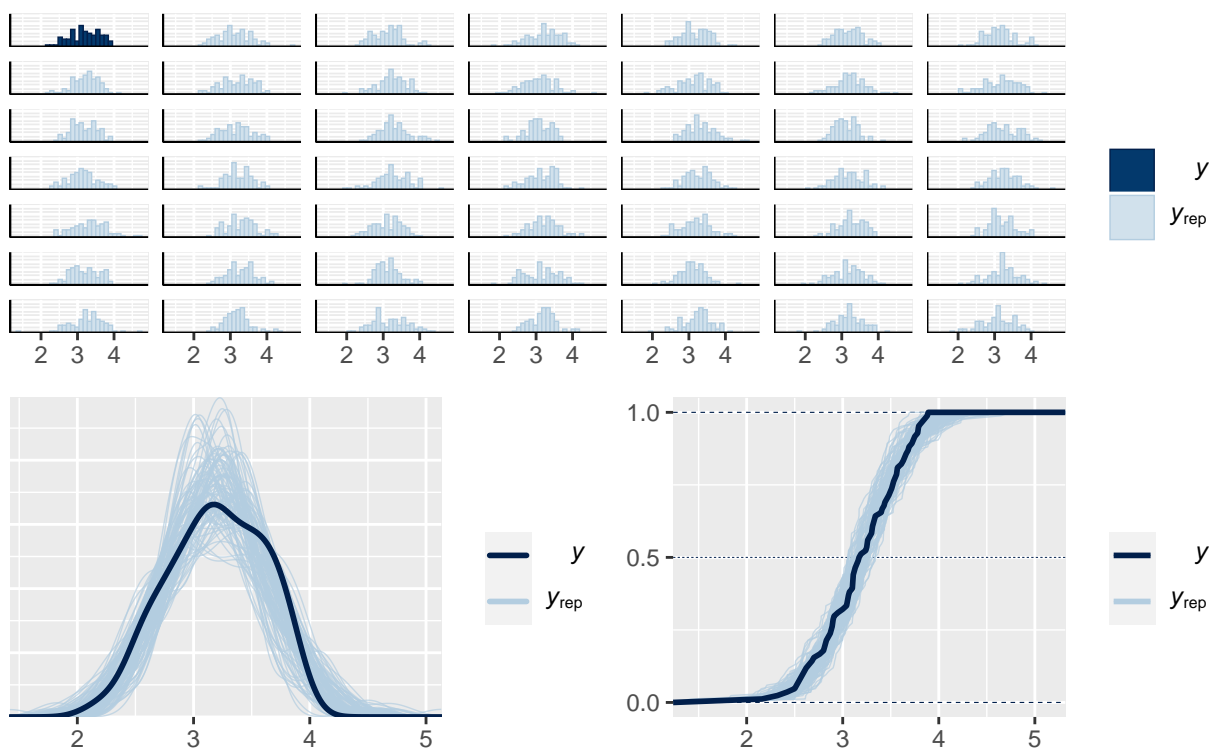


Figure 3: On the top of the figure 48 replications of the hierarchical model for the group of white male students are shown with histograms. The original data is shown on the top left corner. Below the histograms, 100 epdf and ecdf distributions of the replications are overlapped with the epdf and ecdf of original data.

```
y2 <- data_hierarchical$y2
ypred2 <- fit_hierarchical$draws("ypred2", format = "matrix")

grid.arrange(ppc_hist(y2, ypred2[1:48,]),
ppc_dens_overlay(y2, ypred2[1:100,]),
ppc_ecdf_overlay(y2, ypred2[1:100,]),
              layout_matrix = matrix(c(1,2,1,3), nrow = 2), nrow=2)
```

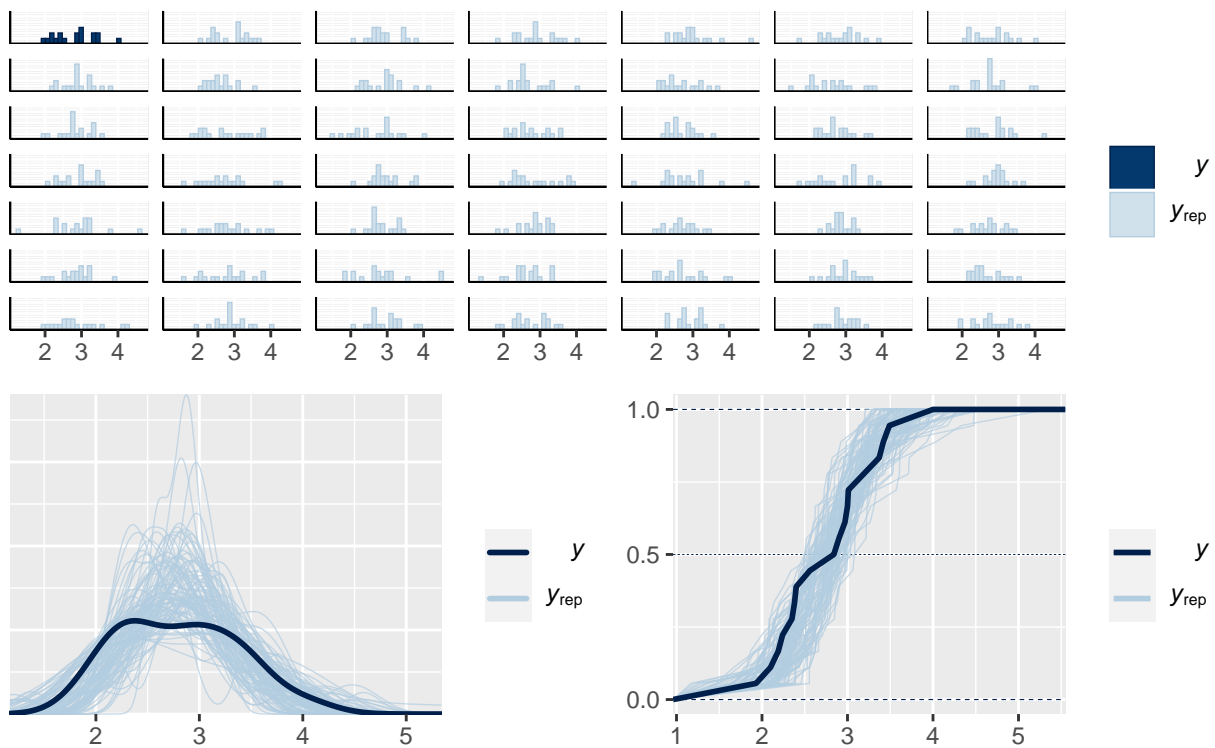


Figure 4: On the top of the figure 48 replications of the hierarchical model for the group of non-white male students are shown with histograms. The original data is shown on the top left corner. Below the histograms, 100 epdf and ecdf distributions of the replications are overlapped with the epdf and ecdf of original data.

```

y3 <- data_hierarchical$y3
ypred3 <- fit_hierarchical$draws("ypred3", format = "matrix")

grid.arrange(ppc_hist(y3, ypred3[1:48,]),
ppc_dens_overlay(y3, ypred3[1:100,]),
ppc_ecdf_overlay(y3, ypred3[1:100,]),
              layout_matrix = matrix(c(1,2,1,3), nrow = 2), nrow=2)

```

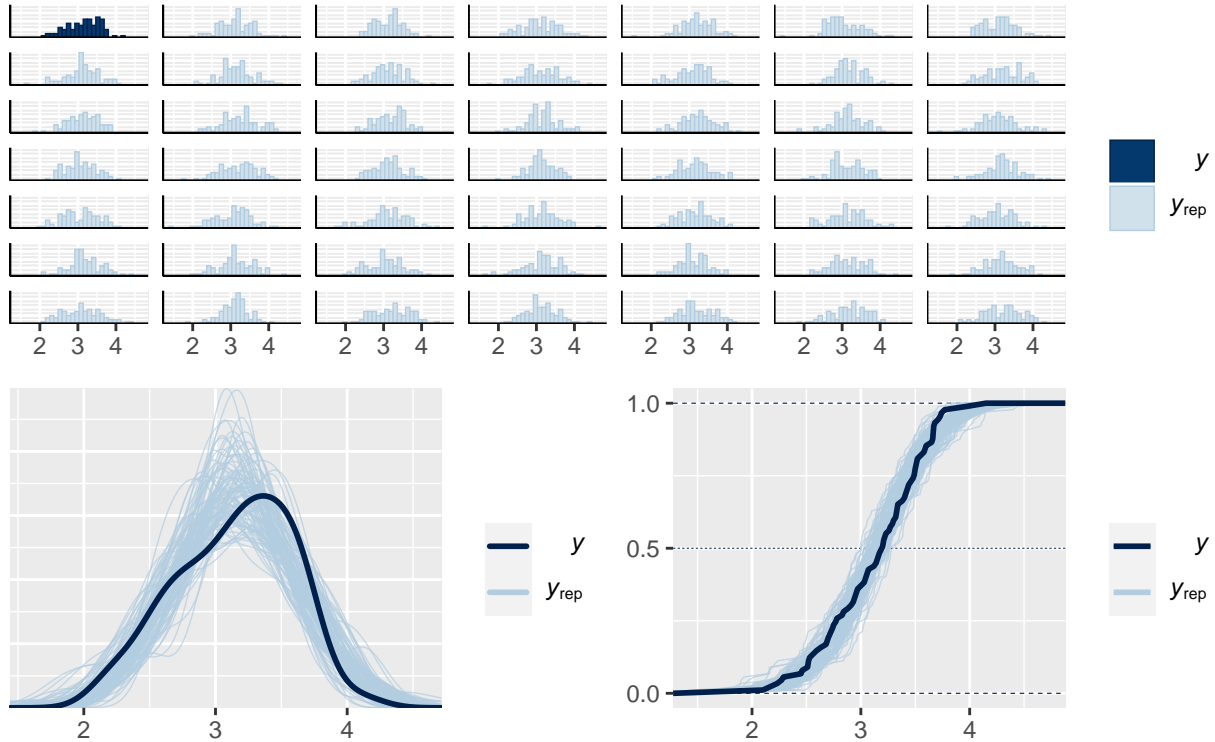


Figure 5: On the top of the figure 48 replications of the hierarchical model for the group of white female students are shown with histograms. The original data is shown on the top left corner. Below the histograms, 100 epdf and ecdf distributions of the replications are overlapped with the epdf and ecdf of original data.

```

y4 <- data_hierarchical$y4
ypred4 <- fit_hierarchical$draws("ypred4", format = "matrix")

grid.arrange(ppc_hist(y4, ypred4[1:48,]),
ppc_dens_overlay(y4, ypred4[1:100,]),
ppc_ecdf_overlay(y4, ypred4[1:100,]),
              layout_matrix = matrix(c(1,2,1,3), nrow = 2), nrow=2)

```




Figure 6: On the top of the figure 48 replications of the hierarchical model for the group of non-white female students are shown with histograms. The original data is shown on the top left corner. Below the histograms, 100 epdf and ecdf distributions of the replications are overlapped with the epdf and ecdf of original data.

Table 1: The priors utilised in prior sensitivity analysis.

alpha, beta OR mu_alpha, mu_beta	sigma, sigma_alpha, sigma_beta
N(0,100)	N(0,10)
N(0,1000)	N(0,10)
N(0,50)	N(0,10)
N(0,10)	N(0,10)
N(0,1)	N(0,10)
N(0,100)	N(0,100)
N(0,100)	N(0,50)
N(0,100)	N(0,1)

Table 2: Sensitivity analysis of the pooled model: Posterior means

	1	2	3	4	5	6	7	8
alpha	0.3901930	NA	NA	NA	NA	NA	NA	NA
beta1	0.4611466	NA	NA	NA	NA	NA	NA	NA
beta2	0.0009451	NA	NA	NA	NA	NA	NA	NA
beta3	0.0003677	NA	NA	NA	NA	NA	NA	NA
beta4	0.0184211	NA	NA	NA	NA	NA	NA	NA
beta5	0.0091861	NA	NA	NA	NA	NA	NA	NA

4.3 Predictive performance assessment

Predictive performance assessment was not performed for the models. The reason for this is that the models are regression models, and no sensible metrics for examining the performance of regression models exist. Theoretically, a metric such as mean absolute error (MAE) or mean squared error (MSE) between the real and predicted values could be calculated for the data. This is not common practice, however, so predictive performance assessment is simply skipped here.

4.4 Prior sensitivity analysis

The original priors utilised were $N(0, 100)$ for α and β or their means, and $N(0, 10)$ for all standard deviations. For prior sensitivity analysis, the priors presented in the Table 1 were tested to see how much the posteriors of α and β_k would differ from the original. Altogether, 8 prior combinations were tested, including the original priors for reference. The same prior combinations were used for testing both models. Narrower priors, e.g. $N(0, 0.1)$, were excluded from the analysis, as they would not be weakly informative anymore.

4.4.1 Pooled model

The pooled model was run with the 8 prior combinations. The resulting posterior means and SDs for all parameters are reported in Tables 2 and 3. The results show that regardless of the used priors, all posteriors for α and β_k are very similar to each other. Therefore it can be concluded that the pooled model is not sensitive to changes in the priors.

4.4.2 Hierarchical model

The hierarchical model’s sensitivity was tested similarly to the pooled model. The resulting posteriors are shown in Tables 4 and 5 only partially, as the full tables would be too voluminous. Furthermore, the results are quite similar for all parameters. There is slightly more variation in the posteriors of the hierarchical model compared to the pooled model when various priors are utilised, but the posteriors look quite similar nonetheless. It can be concluded that the hierarchical model, similarly to the pooled model, is not sensitive to changes in the priors.

Table 3: Sensitivity analysis of the pooled model: Posterior SDs

	1	2	3	4	5	6	7	8
alpha	0.3270551	NA	NA	NA	NA	NA	NA	NA
beta1	0.0742112	NA	NA	NA	NA	NA	NA	NA
beta2	0.0003850	NA	NA	NA	NA	NA	NA	NA
beta3	0.0004237	NA	NA	NA	NA	NA	NA	NA
beta4	0.0039195	NA	NA	NA	NA	NA	NA	NA
beta5	0.0056140	NA	NA	NA	NA	NA	NA	NA

Table 4: Sensitivity analysis of the hierarchical model: Posterior means

	1	2	3	4	5	6	7	8
alpha1	0.6750332	NA	NA	NA	NA	NA	NA	NA
alpha2	0.4111187	NA	NA	NA	NA	NA	NA	NA
alpha3	0.7128101	NA	NA	NA	NA	NA	NA	NA
alpha4	0.3897968	NA	NA	NA	NA	NA	NA	NA
beta1_1	0.4057357	NA	NA	NA	NA	NA	NA	NA
beta2_1	0.0009161	NA	NA	NA	NA	NA	NA	NA
beta3_1	0.0004205	NA	NA	NA	NA	NA	NA	NA
beta4_1	0.0164358	NA	NA	NA	NA	NA	NA	NA
beta5_1	0.0047366	NA	NA	NA	NA	NA	NA	NA

Table 5: Sensitivity analysis of the hierarchical model: Posterior SDs

	1	2	3	4	5	6	7	8
alpha1	0.4894600	NA	NA	NA	NA	NA	NA	NA
alpha2	0.5651351	NA	NA	NA	NA	NA	NA	NA
alpha3	0.4263073	NA	NA	NA	NA	NA	NA	NA
alpha4	0.4793913	NA	NA	NA	NA	NA	NA	NA
beta1_1	0.1048547	NA	NA	NA	NA	NA	NA	NA
beta2_1	0.0005564	NA	NA	NA	NA	NA	NA	NA
beta3_1	0.0006127	NA	NA	NA	NA	NA	NA	NA
beta4_1	0.0054367	NA	NA	NA	NA	NA	NA	NA
beta5_1	0.0078472	NA	NA	NA	NA	NA	NA	NA

Table 6: Posterior distributions of the parameters for both models

	pooled	h. group 1	h. group 2	h. group 3	h. group 4
alpha	N(0.39, 0.327)	N(0.675, 0.489)	N(0.411, 0.565)	N(0.713, 0.426)	N(0.39, 0.479)
beta1	N(0.461, 0.074)	N(0.406, 0.105)	N(0.46, 0.13)	N(0.522, 0.094)	N(0.555, 0.142)
beta2	N(0.001, 0)	N(0.001, 0.001)	N(0, 0.001)	N(0.001, 0.001)	N(0, 0.001)
beta3	N(0, 0)	N(0, 0.001)	N(0.001, 0.001)	N(0, 0.001)	N(0, 0.001)
beta4	N(0.018, 0.004)	N(0.016, 0.005)	N(0.023, 0.009)	N(0.017, 0.005)	N(0.017, 0.008)
beta5	N(0.009, 0.006)	N(0.005, 0.008)	N(0.027, 0.018)	N(0.005, 0.008)	N(0.007, 0.013)

4.5 Posterior distributions

Next, the posterior distributions of α and β_k are inspected and visualised. Table 6 contains the posterior distributions of the six variables for the pooled model and all 4 groups of the hierarchical model separately. All parameters are reported with the accuracy of 3 decimals.

4.5.1 Pooled model

Figure 7 shows that the slope parameter β_1 , multiplying the variable for high school GPA, likely differs from 0 and it the largest of the coefficients β_k . This implies that high school GPA has the most effect on college GPA, our target variable. The other β s have very narrow distributions close to 0, implying their effect on college GPA is either very small or non-existent.

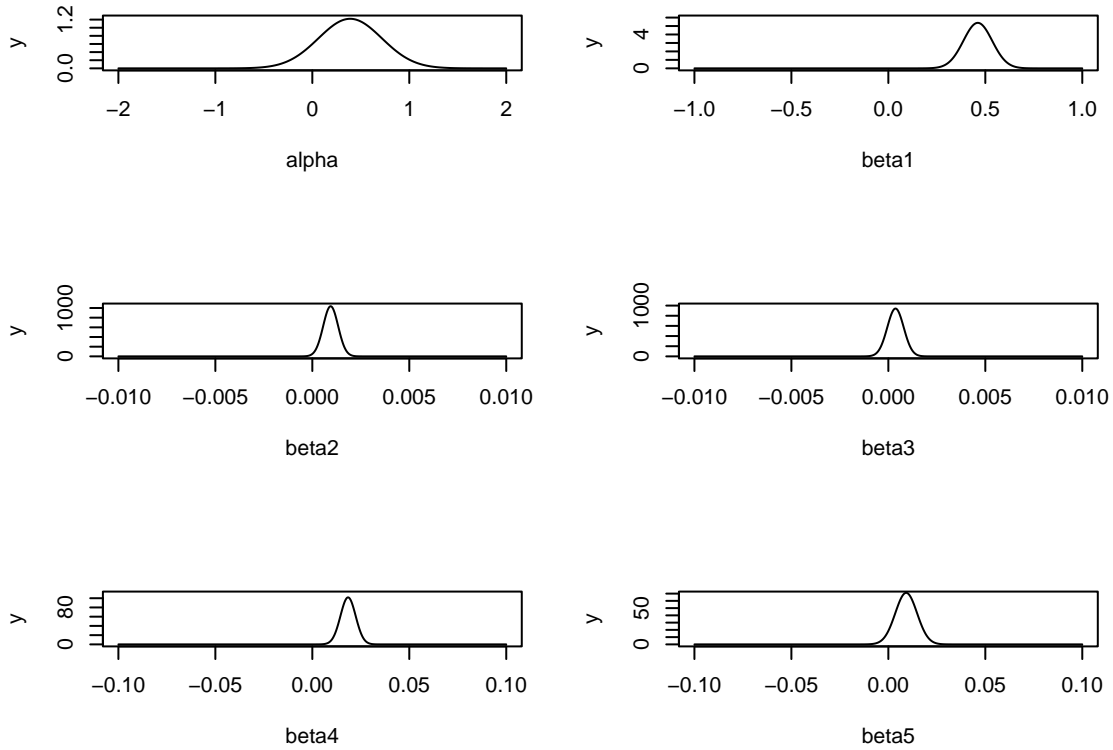


Figure 7: Posterior distributions of the parameters of the pooled model. Note the varying x-axes.

4.5.2 Hierarchical model

From Figure 8, it can be seen that posteriors of the hierarchical model are quite similar to the pooled model. The distributions of the intercept terms α of groups 1 and 3, and of groups 2 and 4, are very similar to each other, implying that the ethnicity of the student governs the intercept term α in the hierarchical model. Furthermore, there is slight variation in the slope terms β_k , especially β_4 and β_5 : for groups with less data, the distributions are narrower, and for groups with more data, wider.

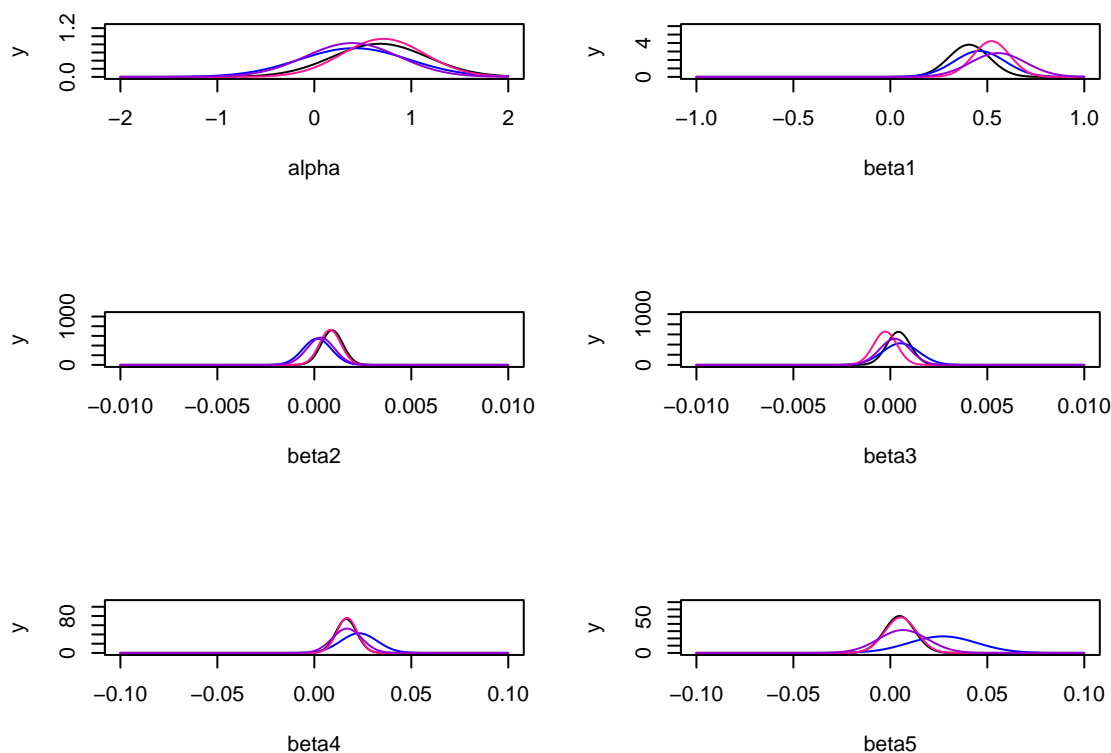


Figure 8: Posterior distributions of the parameters of the hierarchical model. Each group's parameters are plotted in a different colour (black = group 1, blue = group 2, pink = group 3, violet = group 4). Note the varying x-axes.

4.6 Model comparison

The log likelihoods of the data have been calculated in the “generated quantities” blocks of the models. They are given to the loo function of the “loo” package, and the individual results and model comparison are printed.

```
log_lik_pooled <- fit_pooled$draws("log_lik", format = "matrix")
log_lik_hierarchical <- fit_hierarchical$draws("log_lik", format = "matrix")

loo_pooled <- loo(log_lik_pooled)
loo_hierarchical <- loo(log_lik_hierarchical)

loo_pooled
```

```
##
```

```
## Computed from 4000 by 219 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo   -108.0  9.5
## p_loo       6.9   0.7
## looic       216.0 19.0
## -----
## Monte Carlo SE of elpd_loo is 0.0.
##
## All Pareto k estimates are good (k < 0.5).
## See help('pareto-k-diagnostic') for details.
```

```
loo_hierarchical
```

```
##
## Computed from 4000 by 219 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo   -109.7  9.5
## p_loo       18.4  2.4
## looic       219.4 19.0
## -----
## Monte Carlo SE of elpd_loo is NA.
##
## Pareto k diagnostic values:
##           Count Pct.   Min. n_eff
## (-Inf, 0.5] (good)   214  97.7%   877
## (0.5, 0.7]  (ok)      4    1.8%   900
## (0.7, 1]    (bad)      1    0.5%    21
## (1, Inf)    (very bad) 0    0.0%   <NA>
## See help('pareto-k-diagnostic') for details.
```

```
loo_compare(list("pooled" = loo_pooled, "hierarchical" = loo_hierarchical))
```

```
##           elpd_diff se_diff
## pooled           0.0     0.0
## hierarchical    -1.7    14.0
```

The results show that the PSIS-LOO estimate for the pooled model can be considered reliable, as all values $\hat{k} \lesssim 0.7$. The same cannot be said of the hierarchical model, and the PSIS-LOO estimate for it should be considered, or at least suspected, biased and overly optimistic.

Comparing the two models reveals that the pooled model seems to fit to the data better than the hierarchical model (likely due to the lack of training data for the latter one). Therefore for any future predictions, the pooled model should be utilised.

5 Discussion about problems and improvements

The largest problem in the project was the amount of data. Although there was enough data for the pooled model, for the hierarchical model some of the data groups were quite small. For example, the non-white groups only contained 18 and 28 data points, and the rest of the 173 data points were distributed between white groups.

There was large variation in the magnitudes of the variables: for example, the SATV and SATM variables had 100 times larger values than the HSGPA. Because of the magnitude difference, interpreting the coefficients β_k was not very simple. This problem could have been solved by using variable normalisation before running the models.

There were some divergences in the hierarchical model. If there had been more time, the divergences could have been solved by analysing the model and how it was run further. For example, the default values (number of chains, number of iterations etc.) were utilised during model fitting, and different values could be tested. Furthermore, different priors could have been tested to see if it would solve the divergence issue.

6 Conclusion

Both models predicted that the high school GPA affected our target variable, college GPA, the most, as its coefficient β_1 was the largest. Other variables had either a smaller or non-existent effect. It is difficult to say for sure, as the posteriors of some of the β_k overlap with 0, although they are more on the positive side. Furthermore, the large magnitude differences between variables make interpreting the coefficients difficult. The hierarchical model found that the ethnicity of a student governs the α and β_2 parameters and that gender seemed to have no effect on them. However, the parameters were quite similar across all four groups in the hierarchical model, except for a few cases when there was a large difference between the amount of data in the groups. This can be seen e.g. from β_5 : the second group had the least data, and therefore the posterior of β_5 for the second group is the widest.

Based on the model comparison with LOO-CV, the pooled model describes the data better than the hierarchical model. However, both models performed quite well and the performance difference was not large. As there were various issues with the hierarchical model, e.g. divergences and \hat{k} values, the pooled model should be chosen as the better model and utilised for any future predictions.

7 Self-reflection

During the project, we learned to make more complicated Stan models. Additionally, we noticed that finding suitable data for data analysis is surprisingly difficult and time-consuming. Moreover, understanding the data, pre-processing it, and deciding the suitable groups for the hierarchical model took a considerable amount of time. We also noticed the importance of the amount of data: the more data we can feed the model, the better the model fits and predicts. However, we also learned that too much data can become a practical issue, as running the models with all 219 data points took quite a long time.